



On-Device AI Assistant – ASCII Workflow Map, UI Wireframes & Product Spec (MVP)


1. ASCII WORKFLOW MAP (MVP)





2. UI WIREFRAMES (ASCII)

2.1 Chat Interface Wireframe



[Mic]

Type your message here...

[Send]

2.2 Document Ingestion Wireframe

Upload Documents

[Drag & Drop Files Here]

File Queue:

- client_overview.pdf [#####-----] 52%
- portfolio_2024.docx [Processing...]
- archive.xls [Failed: Unsupported]

Status: Embedding documents locally...

2.3 Document Preview Panel Wireframe

Chat Window

AI: Here's your summary...
[1][2]

User: Expand on risk...

Sources (3)

- portfolio_2024.docx
Pg 2: "Market vol..."
- client_overview.pdf
Pg 5: "Liquidity..."

2.4 Settings Wireframe

Settings

Storage Location: [/Users/data/assistant/] (Change)

Enable Encryption: [X] Yes

Enable Audit Logs: [X] Yes

Clear Local Data: [Clear All]

Dark Mode:

[X] Enabled

2.5 Audit Log Viewer Wireframe

Audit Logs		
Timestamp	Query	Docs Used
10:04 AM	"Summarize client risk..."	client_overview
10:06 AM	"Explain allocation..."	portfolio_2024
[Search Logs...] [Export Logs]		

3. PRODUCT SPECIFICATION (MVP)

3.1 Product Overview

A fully offline AI assistant for Private Banking Wealth Advisors. Provides secure document-grounded answers, internal research support, and fast retrieval of client-relevant insights.

3.2 Goals & Non-Goals

Goals

- Offline LLM interaction.
- Secure document ingestion & storage.
- Fast retrieval-augmented answers.
- Transparent sources for compliance.
- Exportable insights.

Non-Goals

- No cloud connectivity.
- No real-time market feeds.
- No client-facing UI.

3.3 Users

Primary: Private Banking Wealth Advisors.

3.4 Core Features (MVP)

1. **Chat Interface** – real-time Q&A with local LLM.
 2. **Document Ingestion** – PDF/Doc ingestion → chunking → embeddings.
 3. **RAG Pipeline** – local vector search + answer generation.
 4. **Document Source Panel** – citation-aware UI.
 5. **Settings** – encryption, storage, logs.
 6. **Audit Logs** – transparency + compliance.
 7. **System Notifications** – ingestion progress, errors.
-

3.5 Functional Requirements

Chat

- User can submit queries.
- LLM must generate answers using local document store.
- Support markdown formatting.
- Display typing indicator.

Document Ingestion

- Upload via file picker or drag-drop.
- Extract → Chunk → Embed locally.
- Show per-file progress.
- Allow re-ingesting updated docs.

Search & Sources

- Show ranked document chunks.
- Open preview panel.
- Highlight matched terms.

Settings

- Change storage directory.
- Enable encryption at rest.
- Toggle audit logging.

Audit Logs

- Log query text, timestamp, and referenced docs.
 - Export logs to CSV or text.
-

3.6 Technical Requirements

- Electron app environment.
- Local LLM (e.g., GGUF model via llama.cpp).
- Local vector DB (e.g., SQLite + Faiss / Milvus Lite).

- All processing completely offline.
 - Data encrypted at rest (AES-256 optional).
-

3.7 UX / Design Principles

- Simple, secure, advisor-focused.
 - Vision Pro-inspired translucent, clean design.
 - Accessible typography + high contrast.
 - No clutter; emphasize insight.
-

3.8 Edge Cases

- Empty document store → Chat disabled.
 - Corrupted file → Error banner.
 - Low disk space → Warning modal.
-

3.9 Future Enhancements (Post-MVP)

- Voice input.
 - Light mode.
 - Collaborative sharing of local knowledge base.
 - Multi-model selection.
 - Automated scheduled re-indexing.
-

END OF DOCUMENT