

Product Requirements Document (PRD)

Wealth Companion: AI-Powered Desktop Widget for Wealth Management Professionals

Document Information

Product Name: Wealth Companion
Version: 1.0 (MVP)
Status: Draft for Review
Target Release: Q2 2026
Document Owner: Product Management
Last Updated: October 18, 2025
Classification: Internal - Confidential

Executive Summary

Wealth Companion is an AI-powered desktop widget designed to transform how wealth management professionals—advisors, traders, compliance officers, and relationship managers—work with complex financial systems. By combining ambient intelligence with autonomous computer use capabilities, the product delivers a persistent, context-aware assistant that reduces administrative overhead, enhances client service quality, and enables proactive opportunity identification across client portfolios.

The MVP focuses on a single high-impact use case: **Pre-Meeting Client Intelligence Briefing**, delivering measurable productivity gains while establishing technical foundations for future expansion.

Strategic Value Proposition:

- **50+ minutes saved per advisor per week** through automated research and data aggregation
- **Enhanced client experience** through comprehensive, context-rich meeting preparation
- **Proactive opportunity identification** via continuous portfolio and market monitoring
- **Reduced compliance risk** through automated checks and documentation

1. Strategic Context

1.1 Problem Statement

Wealth management professionals face a critical productivity paradox: they must deliver highly personalized, relationship-driven service while navigating an increasingly complex technology landscape [94][98].

Current Pain Points:

- **Fragmented Systems:** Advisors toggle between 8-12 applications daily (CRM, portfolio management, trading systems, email, calendar, research platforms, compliance tools) [96][98][^101]
- **Manual Research Burden:** Pre-meeting preparation requires 10-15 minutes of manual data gathering across multiple systems [96][113]
- **Context Switching Cost:** Each system transition costs 20-30 minutes of cognitive overhead, destroying focus and productivity [98][101]
- **Missed Opportunities:** Lack of real-time portfolio monitoring means advisors miss trigger events (rebalancing needs, tax-loss harvesting, life event-based opportunities) [80][86][^96]
- **Administrative Overload:** 40-60% of advisor time spent on administrative tasks rather than client-facing activities [80][83][^86]

1.2 Strategic Alignment

This product directly supports organizational objectives:

Business Goals:

- Increase advisor productivity by 15-20% within 12 months [80][86]
- Improve client retention through enhanced service quality [83][98]
- Enable scalable growth without proportional headcount increases [83][86]
- Reduce compliance incidents and associated costs [99][102]

Competitive Differentiation:

- First-mover advantage in agentic AI for wealth management [102][111]
- Proprietary integration with internal systems creates moat [94][98]
- Ambient computing interface differentiates from chatbot-only solutions [1][2][3]

Market Opportunity:

- Global WealthTech market projected to reach \$2.5T by 2030 [80][83]
- 68% of financial institutions prioritizing AI agent investments [102][111]
- Average efficiency gains of 43% reported from AI agent implementations [64][102]

1.3 User Personas

Primary Persona: Client-Facing Financial Advisor

- **Demographics:** 35-55 years old, 5-15 years experience, manages 80-150 UHNW/HNW client relationships
- **Daily Activities:** Client meetings (3-5/day), portfolio reviews, research, prospecting, administrative tasks
- **Pain Points:** Insufficient time for relationship building, manual data gathering, context switching between systems
- **Goals:** Deliver exceptional personalized service, grow AUM, maintain work-life balance
- **Tech Proficiency:** Moderate; comfortable with CRM and portfolio tools but not coding

Secondary Persona: Portfolio Manager/Trader

- **Demographics:** 30-50 years old, quantitative background, manages execution and portfolio construction
- **Daily Activities:** Trade execution, portfolio monitoring, compliance checks, market research
- **Pain Points:** Pre-trade compliance verification, manual order aggregation, risk monitoring
- **Goals:** Optimal execution quality, zero compliance violations, efficient workflows
- **Tech Proficiency:** High; comfortable with terminal systems and APIs

Tertiary Persona: Compliance Officer

- **Demographics:** 35-60 years old, regulatory expertise, risk management focus
- **Daily Activities:** Trade surveillance, policy enforcement, regulatory monitoring, audit preparation
- **Pain Points:** Manual review processes, keeping pace with regulatory changes, audit documentation
- **Goals:** Zero violations, efficient monitoring, proactive risk identification
- **Tech Proficiency:** Moderate; proficient with compliance platforms

1.4 Success Metrics (MVP)

Primary Metrics:

- **Time Savings:** Reduce pre-meeting prep time from 12 min to 2 min (83% reduction)
- **Adoption Rate:** 80% of target users actively using widget within 30 days of deployment
- **User Satisfaction:** NPS $\geq$ 40, CSAT $\geq$ 4.2/5.0 within 60 days
- **Accuracy:** 95%+ accuracy on client data synthesis (validated through advisor feedback)

Secondary Metrics:

- System Response Time: ≤ 5 seconds for briefing generation
- Daily Active Usage: Average 3-5 briefings per advisor per day
- Data Completeness: 90%+ of briefings include all required data sections

Business Impact Metrics (6-month horizon):

- Productivity Gain: 50+ minutes saved per advisor per week
- Client Meeting Quality: Measured through post-meeting client feedback (target: 10% improvement)
- Opportunity Identification: Number of proactive client outreach actions generated
- ROI: 5x return on implementation costs within 12 months

2. Product Scope and Use Cases

2.1 MVP Scope: Pre-Meeting Client Intelligence Briefing

In Scope for MVP:

The widget automatically generates comprehensive client briefings 30 minutes before scheduled meetings, aggregating information from multiple systems into a single, structured summary.

Core Functionality:

1. **Calendar Integration:** Monitor advisor calendar for upcoming client meetings
2. **Data Aggregation:** Pull information from CRM, portfolio management system, email, and news sources
3. **Intelligent Synthesis:** Generate structured briefing with key information and insights
4. **Persistent Display:** Present briefing in widget sidebar for easy reference during meeting
5. **Action Tracking:** Allow advisors to mark follow-up items during/after meeting

Briefing Content Sections:

- **Client Overview:** Key demographics, relationship tenure, total AUM, risk profile
- **Recent Activity:** Last meeting date/notes, recent communications, portfolio transactions (past 30 days)
- **Portfolio Performance:** YTD performance vs. benchmark, asset allocation, concentration risks
- **Alerts & Opportunities:** Rebalancing needs, tax-loss harvesting, upcoming life events, market-relevant news
- **Meeting Context:** Agenda from calendar, outstanding action items from previous meetings
- **Conversation Starters:** 2-3 personalized talking points based on client interests and recent events

2.2 Detailed User Story (MVP)

As a financial advisor

I want to receive automated, comprehensive client briefings before meetings

So that I can deliver highly personalized service without spending time on manual research

Acceptance Criteria:

1. Widget automatically detects client meetings in calendar 30 min in advance
2. System retrieves client data from all connected systems (CRM, portfolio, email) within 5 seconds
3. Briefing displays in sidebar with collapsible sections for easy scanning
4. All data points include source attribution and last-updated timestamps
5. Advisor can mark specific items for follow-up with one click
6. Briefing remains accessible during and after meeting
7. System logs which briefing sections advisor reviewed (for continuous improvement)
8. Advisor can provide feedback rating (1-5 stars) and comments on briefing quality

User Flow:

1. Advisor starts workday; widget icon visible in system tray/menu bar
2. 30 minutes before meeting, widget expands and shows notification: "Client briefing ready for [Client Name]"
3. Advisor clicks notification; widget sidebar expands showing briefing
4. Advisor scans key sections (2 minutes); mentally prepares for meeting
5. During meeting, advisor references briefing on second monitor
6. After meeting, advisor clicks "Mark complete" and optionally adds follow-up tasks
7. Widget contracts back to minimal mode

2.3 Out of Scope for MVP

Explicitly NOT included in initial release:

- Multi-step autonomous actions (e.g., auto-sending emails, executing trades)
- Voice interaction capabilities
- Integration with external systems outside core tech stack
- Mobile application or web interface (desktop-only)
- Proactive opportunity scanning outside meeting context
- Advanced NLP for email/document analysis
- Custom workflow builder
- Team collaboration features
- Multi-language support (English only)
- Historical briefing analytics and trends

Future Roadmap Considerations (Post-MVP):

- **Phase 2 (Q3-Q4 2026):** Prospect research automation, portfolio health monitoring, regulatory change tracking
- **Phase 3 (2027):** Autonomous workflows, voice interaction, team knowledge sharing, mobile companion

3. Functional Requirements

3.1 Core System Components

3.1.1 Widget Interface Layer

FR-UI-001: Persistent Desktop Widget

- Widget appears as dockable sidebar on desktop (Windows 10/11, macOS 12+)
- Three display states:
 - **Minimal:** Small icon in system tray/menu bar (20x20px)
 - **Compact:** Notification card (300x200px) for alerts
 - **Expanded:** Full sidebar panel (400px wide, full screen height)
- Smooth animated transitions between states ($\leq 300\text{ms}$)
- User-configurable docking position (left/right screen edge)
- Always-on-top option for during-meeting reference

FR-UI-002: Briefing Display Interface

- Structured layout with collapsible sections
- Clear visual hierarchy (headings, spacing, typography)
- Inline source attribution for all data points

- Color-coded alerts (green=positive, yellow=attention needed, red=urgent)
- Scrollable content with sticky header showing client name
- Print-to-PDF capability for offline reference

FR-UI-003: Interaction Controls

- "Mark complete" button to dismiss active briefing
- Quick-add follow-up task with text input
- Thumbs up/down feedback mechanism
- Expand/collapse all sections toggle
- "Refresh" button to regenerate briefing with latest data
- Settings gear icon for configuration access

3.1.2 Calendar Integration Module

FR-CAL-001: Meeting Detection

- Monitor advisor's calendar (Outlook, Google Calendar) via API
- Identify client meetings based on:
 - Meeting attendee matching CRM client records
 - Calendar event title containing client name
 - Custom meeting types/categories (configurable)
- Trigger briefing generation 30 minutes before meeting start time
- Handle meeting time changes (update briefing if rescheduled)

FR-CAL-002: Context Extraction

- Parse meeting agenda from calendar event body
- Extract meeting type (annual review, check-in, planning session)
- Identify other attendees (internal colleagues, external professionals)
- Capture meeting location for in-person vs virtual distinction

3.1.3 Data Integration Layer

FR-INT-001: CRM Integration

- Real-time read access to internal CRM via REST API
- Retrieved fields:
 - Client demographics (name, age, occupation, family structure)
 - Relationship details (start date, relationship manager, service tier)
 - Contact information (email, phone, address)
 - Custom fields (interests, life events, preferences)
 - Historical notes from previous meetings
 - Outstanding action items/tasks
- Data refresh on each briefing generation
- Graceful degradation if CRM unavailable (show cached data with warning)

FR-INT-002: Portfolio Management System Integration

- Read access to portfolio platform via API or database query
- Retrieved data:
 - Current holdings (positions, quantities, market values)
 - Performance metrics (YTD, 1-year, 3-year returns)

- Asset allocation breakdown
- Benchmark comparison
- Recent transactions (past 30 days)
- Pending orders
- Risk metrics (beta, standard deviation, Sharpe ratio)
- Calculate derived insights:
 - Drift from target allocation
 - Concentration risks (single position > 10% of portfolio)
 - Tax-loss harvesting opportunities

FR-INT-003: Email System Integration

- Search recent email threads with client (past 90 days)
- Extract key topics and unresolved questions
- Identify attachments shared (proposals, research reports)
- Flag unanswered client emails or pending responses
- Summarize recent correspondence (3-5 sentence summary)

FR-INT-004: Market Data & News Integration

- Monitor news sources for client-relevant mentions (Bloomberg, Reuters, internal research)
- Track sector/industry news relevant to client portfolio holdings
- Flag significant market events affecting portfolio (earnings announcements, regulatory changes)
- Pull latest market data for key positions

3.1.4 AI Processing Engine

FR-AI-001: Data Synthesis

- Aggregate data from all integrated systems into structured format
- Use LLM (GPT-4o, Claude 3.5 Sonnet) to generate briefing narrative
- Apply prompt engineering to ensure consistent output format
- Include confidence scoring for synthesized insights
- Maintain context window limit ($\leq 128K$ tokens input)

FR-AI-002: Insight Generation

- Identify actionable opportunities:
 - Rebalancing triggers (drift > 5% from target)
 - Tax-loss harvesting (positions down > 10% with tax lot context)
 - Life event triggers (birthdays, anniversaries, approaching retirement)
 - Market-driven opportunities (volatility, sector rotation)
- Generate 2-3 personalized conversation starters based on:
 - Client interests/hobbies from CRM notes
 - Recent news about client's company or industry
 - Portfolio performance relative to goals

FR-AI-003: Reasoning and Explanations

- Provide transparent reasoning for insights
- Cite specific data sources for each statement
- Include confidence levels (High/Medium/Low) for synthesized insights
- Allow drill-down to source data

3.1.5 Security & Compliance Module

FR-SEC-001: Authentication & Authorization

- Single Sign-On (SSO) integration with corporate identity provider
- Role-based access control (RBAC):
 - Advisors: Access only their assigned clients
 - Managers: Access team's clients
 - Compliance: Read-only access across organization
- Session management with auto-logout after inactivity (15 min)
- Multi-factor authentication (MFA) enforcement

FR-SEC-002: Data Protection

- End-to-end encryption for data in transit (TLS 1.3)
- Encryption at rest for cached data (AES-256)
- No persistent storage of client data in widget (ephemeral)
- Secure credential management (no hardcoded secrets)
- Data residency compliance (all processing within regional boundaries)

FR-SEC-003: Audit Logging

- Log all briefing generation events (timestamp, user, client, data sources)
- Log user interactions (sections viewed, tasks created, feedback provided)
- Tamper-proof audit trail for compliance review
- Retention period: 7 years (configurable per regulatory requirements)
- Export capability for audit requests

FR-SEC-004: Compliance Controls

- Pre-flight data validation (ensure all required fields present)
- Automated PII detection and masking for logs
- Compliance approval workflow for initial deployment
- Regular security scanning and penetration testing
- Incident response protocol

3.2 Non-Functional Requirements

3.2.1 Performance Requirements

NFR-PERF-001: Response Time

- Briefing generation: ≤ 5 seconds (95th percentile)
- Widget state transition: ≤ 300 ms
- Data refresh: ≤ 3 seconds
- UI interactions: ≤ 100 ms feedback

NFR-PERF-002: Scalability

- Support 500 concurrent users (initial deployment)
- Plan for 2,000 users within 18 months
- Handle 50 briefing requests per minute during peak hours
- Graceful degradation under load (queue requests, show wait times)

NFR-PERF-003: Reliability

- System uptime: 99.5% (excluding scheduled maintenance)

- Maximum briefing generation failure rate: < 1%
- Automatic retry mechanism for transient failures (3 attempts with exponential backoff)
- Circuit breaker pattern for dependent services

3.2.2 Usability Requirements

NFR-USE-001: Ease of Use

- Zero-touch onboarding: Widget auto-installs and configures via enterprise software deployment
- First briefing generated within 2 minutes of installation
- Maximum 3 clicks to access any feature
- Inline contextual help (tooltips, help icons)
- Comprehensive user guide and video tutorials

NFR-USE-002: Accessibility

- WCAG 2.1 Level AA compliance
- Keyboard navigation support
- Screen reader compatibility
- Adjustable font sizes (3 presets: Small, Medium, Large)
- High-contrast mode option

3.2.3 Compatibility Requirements

NFR-COMP-001: Desktop Platforms

- Windows 10 (version 1909+) and Windows 11
- macOS 12 (Monterey) and later
- Minimum screen resolution: 1920x1080 (for optimal experience)
- Multi-monitor support (widget on secondary display)

NFR-COMP-002: System Requirements

- Minimum 8GB RAM (16GB recommended)
- 500MB available disk space
- Active internet connection (100 Mbps recommended)
- Supported browsers for web components: Chrome 90+, Edge 90+, Safari 14+

3.2.4 Data & Integration Requirements

NFR-DATA-001: Data Freshness

- CRM data: Real-time (on-demand refresh)
- Portfolio data: Real-time during market hours, T+1 for end-of-day valuations
- Email data: Past 90 days, refreshed on briefing generation
- News data: Real-time via streaming APIs

NFR-DATA-002: Data Quality

- Data completeness: 95%+ of required fields populated
- Automated data quality checks before briefing generation
- Clear indication of missing or stale data in briefing
- Fallback logic when primary data sources unavailable

4. User Experience & Design

4.1 Design Principles

1. **Ambient Intelligence:** Proactive assistance without intrusive interruptions
2. **Information Density:** Maximum insight in minimum space
3. **Scannable Content:** Visual hierarchy enables 30-second comprehension
4. **Transparency:** Always show data sources and last-updated timestamps
5. **Minimalist Aesthetics:** Clean, professional interface aligned with enterprise standards

4.2 Visual Design Guidelines

Color Palette:

- Primary: Corporate brand blue (#0052CC)
- Success: Green (#00875A) - positive performance, completed items
- Warning: Amber (#FFAB00) - attention needed
- Error: Red (#DE350B) - urgent items, risks
- Neutral: Grays (#091E42, #344563, #97A0AF, #F4F5F7)

Typography:

- Primary Font: Inter (system font fallback: -apple-system, system-ui)
- Heading Sizes: H1 (18px), H2 (16px), H3 (14px)
- Body: 13px regular, 13px medium for emphasis
- Monospace: SF Mono (code, data values)

Spacing:

- Base unit: 8px (all spacing in 8px increments)
- Section padding: 16px
- Element margins: 8px
- Card radius: 8px

4.3 Key Screens/Views

View 1: Minimal State (Resting)

- Small icon in system tray/menu bar
- Badge indicator shows number of pending briefings
- Hover shows tooltip: "Wealth Companion - Next briefing in 25 min"

View 2: Notification State (30 min before meeting)

- Compact card appears in screen corner (notification style)
- Content:
 - Client name and photo (if available)
 - Meeting time and type
 - "View Briefing" button
- Auto-dismisses after 60 seconds if not clicked (but briefing remains accessible)

View 3: Expanded Briefing View (Primary Interface)

- Full sidebar (400px wide)
- Sticky header:
 - Client name and photo

- Meeting time and countdown
- Quick actions: Mark complete, Refresh, Settings
- Scrollable body with sections:
 1. Client Overview (always visible)
 2. Recent Activity
 3. Portfolio Snapshot
 4. Alerts & Opportunities (highlighted)
 5. Conversation Starters
 6. Meeting Context
- Footer: Feedback mechanism ("Was this helpful?")

View 4: Settings/Configuration

- Overlay panel within widget
- Options:
 - Notification timing (15/30/60 min before meetings)
 - Display preferences (compact/detailed)
 - Data sources to include/exclude
 - Docking position
 - Theme (light/dark)
- Save and Cancel buttons

4.4 Interaction Patterns

Pattern 1: Progressive Disclosure

- Sections collapsed by default except "Alerts & Opportunities"
- Click section header to expand/collapse
- "Expand All" / "Collapse All" toggle at top

Pattern 2: Inline Actions

- Hover over opportunity insight reveals "Add to Follow-up" button
- Click adds item to advisor's task list in CRM
- Visual confirmation (checkmark animation)

Pattern 3: Source Attribution

- Info icon (i) next to data points
- Hover shows tooltip: "Source: CRM | Last updated: 2h ago"
- Click opens source system to relevant record (deep link)

4.5 Wireframes & Prototypes

(Note: Actual wireframes to be created in Figma and linked here. Below is a text description.)

Briefing Layout Structure:

```
+-----+
| [Client Photo] John Smith          |
| Meeting: 10:30 AM (in 25 min)      |
| Annual Review                      |
| [Mark Complete] [Refresh] [⌘]     |
+-----+
|                                     |
| ▢ CLIENT OVERVIEW [v]             |
| Age: 62 | AUM: $4.5M | Risk: Moderate |
```


```
| Relationship: 8 years |
|
| ▢ ALERTS & OPPORTUNITIES [^] |
| ⚠ Portfolio drift: 7% from target |
|   → [Add to Follow-up] |
| ✔ Tax-loss harvesting: $18K savings |
|   → [Add to Follow-up] |
| ▢ Birthday next week (May 15) |
|
| ▢ PORTFOLIO SNAPSHOT [v] |
| YTD Performance: +8.2% (vs S&P +7.1%) |
| Top Holdings: AAPL (12%), MSFT (9%) |
|
| ▢ CONVERSATION STARTERS [v] |
| • Congrats on daughter's college... |
| • Tech concentration risk discussion |
| • Estate planning with upcoming... |
|
| ▢ MEETING CONTEXT [v] |
| Agenda: Q1 review, discuss 2025... |
| Outstanding: Follow up on Roth... |
|
+-----+
| Was this helpful? ▢ ▢ |
+-----+
```


5. Technical Architecture

5.1 System Architecture Overview

Architecture Pattern: Microservices with event-driven communication

Component Diagram:





5.2 Technology Stack

Frontend (Desktop Widget):

- Framework: Electron 28+ (cross-platform desktop app)
- UI Library: React 18+ with TypeScript
- State Management: Redux Toolkit
- Styling: Tailwind CSS + CSS Modules
- Local Storage: IndexedDB (for caching)
- HTTP Client: Axios with retry logic
- Real-time: Socket.IO client for push notifications

Backend Services:

- Language: Node.js 20+ (primary), Python 3.11+ (AI service)
- API Framework: Express.js (Node) / FastAPI (Python)
- Authentication: OAuth 2.0 / SAML 2.0 (SSO)
- API Gateway: Kong / AWS API Gateway
- Message Queue: RabbitMQ / Apache Kafka
- Task Scheduling: Bull Queue (Redis-backed)

AI/ML Components:

- LLM Providers: OpenAI (GPT-4o), Anthropic (Claude 3.5 Sonnet)
- Model Router: LiteLLM (intelligent routing between providers)
- Vector Database: Pinecone / Weaviate (for future semantic search)
- Prompt Management: LangChain / Custom prompt templates

Data & Storage:

- Primary Database: PostgreSQL 15+ (relational data, audit logs)
- Cache Layer: Redis 7+ (session management, rate limiting)
- Document Store: MongoDB (unstructured logs, temporary data)
- Object Storage: AWS S3 / Azure Blob (documents, backups)

Infrastructure:

- Deployment: Docker containers orchestrated by Kubernetes
- Cloud Provider: AWS / Azure / On-premises (hybrid option)
- CI/CD: GitHub Actions / GitLab CI
- Monitoring: Datadog / New Relic / Prometheus + Grafana
- Logging: ELK Stack (Elasticsearch, Logstash, Kibana)
- Secret Management: HashiCorp Vault / AWS Secrets Manager

5.3 Data Flow: Briefing Generation

Step-by-Step Flow:

1. **Trigger:** Calendar Service detects upcoming meeting (30 min before)
2. **Event:** Publish `BriefingRequested` event to message queue
 - Payload: `{advisorId, clientId, meetingId, meetingTime}`
3. **Orchestration:** Briefing Generation Service consumes event
4. **Data Retrieval (Parallel):**
 - Call CRM Service API → Client demographics, notes, tasks
 - Call Portfolio Service API → Holdings, performance, transactions
 - Call Email Service API → Recent correspondence summary
 - Call News Service API → Client-relevant news
5. **Data Aggregation:** Compile all data into structured JSON
6. **AI Processing:**
 - Send aggregated data + prompt template to AI Service
 - AI Service calls LLM API (GPT-4o / Claude via model router)
 - LLM generates briefing sections (JSON format)
7. **Post-Processing:**
 - Validate output format and completeness
 - Add confidence scores and source citations
 - Store briefing in cache (Redis) with 24-hour TTL
8. **Notification:** Send push notification to Desktop Widget via WebSocket
9. **Display:** Widget receives briefing, renders UI, logs view event

Timing Budget:

- Data retrieval (parallel): 2 seconds
- AI generation: 2.5 seconds
- Post-processing & delivery: 0.5 seconds
- **Total: 5 seconds**

5.4 Integration Specifications

CRM Integration (Internal System):

- Protocol: REST API over HTTPS
- Authentication: OAuth 2.0 service account
- Endpoints:
 - GET `/api/v1/clients/{clientId}` - Full client profile
 - GET `/api/v1/clients/{clientId}/notes` - Recent meeting notes
 - GET `/api/v1/clients/{clientId}/tasks` - Outstanding action items
 - POST `/api/v1/tasks` - Create follow-up task
- Rate Limits: 100 requests/min per service account
- Data Format: JSON

Portfolio Management Integration:

- Protocol: REST API over HTTPS or Direct Database Query (read-only replica)
- Authentication: API key + IP whitelist (if API) / Service account (if DB)
- Endpoints / Queries:

- GET /api/v1/portfolios/{portfolioId} - Current holdings
- GET /api/v1/portfolios/{portfolioId}/performance - Returns/metrics
- GET /api/v1/portfolios/{portfolioId}/transactions?since={date}
- Rate Limits: 200 requests/min
- Data Format: JSON

Email Integration (Microsoft 365):

- Protocol: Microsoft Graph API
- Authentication: OAuth 2.0 with delegated permissions
- Endpoints:
 - GET /v1.0/users/{userId}/messages?\$filter=... - Search emails
 - GET /v1.0/users/{userId}/mailFolders/sentItems/messages - Sent items
- Rate Limits: Per Microsoft Graph throttling limits
- Data Format: JSON

Calendar Integration (Microsoft 365 / Google):

- Protocol: Microsoft Graph API / Google Calendar API
- Authentication: OAuth 2.0
- Endpoints:
 - GET /v1.0/users/{userId}/calendar/events (Microsoft)
 - GET /calendar/v3/calendars/primary/events (Google)
- Webhooks: Subscribe to calendar change notifications
- Rate Limits: Per provider policies
- Data Format: JSON

News & Market Data:

- Providers: Bloomberg API, Reuters, Internal Research Platform
- Protocol: REST API / WebSocket (for streaming)
- Authentication: API keys
- Rate Limits: Per provider SLA

5.5 Security Architecture

Defense in Depth:

1. **Perimeter:** Web Application Firewall (WAF), DDoS protection
2. **Authentication:** SSO with MFA enforcement, JWT tokens (15 min expiry)
3. **Authorization:** RBAC with least-privilege principle
4. **Data Protection:** TLS 1.3 in transit, AES-256 at rest
5. **Application:** Input validation, output encoding, SQL injection prevention
6. **Monitoring:** Real-time threat detection, anomaly detection

Data Privacy:

- Zero persistent client data storage in widget (ephemeral only)
- All cached data encrypted locally
- Automatic cache purge on logout or 24-hour TTL
- PII tokenization where possible
- GDPR compliance: Right to erasure, data portability

Compliance:

- SOC 2 Type II certification (infrastructure)
- Regular penetration testing (quarterly)
- Vulnerability scanning (continuous)
- Security incident response plan (SIRP)
- Compliance audit trail (7-year retention)

5.6 Error Handling & Resilience

Failure Scenarios:

Failure	Detection	Recovery Strategy
CRM API unavailable	Health check timeout (5s)	Use cached data, show staleness warning
LLM API error	HTTP 500/503 from provider	Retry with exponential backoff (3 attempts), failover to alternate provider
Network connectivity lost	Connection timeout	Queue briefing request, retry when connection restored
Invalid data format	Schema validation failure	Log error, show partial briefing with missing sections flagged
Widget crash	Electron crash reporter	Auto-restart widget, restore state from cache

Circuit Breaker Pattern:

- Threshold: 5 failures in 60 seconds
- Open state duration: 30 seconds
- Half-open state: Allow 1 request to test recovery

Graceful Degradation:

- If AI service unavailable: Show raw data without synthesis
- If portfolio data missing: Show CRM data only with note
- If all systems down: Display cached briefing with prominent staleness warning

6. AI/ML Specifications

6.1 Model Selection & Routing

Primary LLM: OpenAI GPT-4o (128K context)

- **Use Cases:** Complex synthesis, multi-source data integration, insight generation
- **Rationale:** Best-in-class reasoning, reliable JSON output, strong financial domain knowledge

Secondary LLM: Anthropic Claude 3.5 Sonnet (200K context)

- **Use Cases:** Fallback for GPT-4o, longer context needs, specific compliance-sensitive tasks
- **Rationale:** Strong reasoning, excellent at following instructions, competitive pricing

Model Router Logic:

- Default: GPT-4o
- Fallback: Claude 3.5 Sonnet (if GPT-4o unavailable or rate limited)
- Context size > 100K tokens: Claude 3.5 Sonnet
- Cost optimization mode: Use cheaper models for simple briefings

6.2 Prompt Engineering

System Prompt Template:

You are an AI assistant for wealth management professionals at [Company Name]. Your role is to generate content for client briefings.

INSTRUCTIONS:

1. Synthesize provided data into a structured briefing format
2. Highlight actionable opportunities and alerts
3. Generate personalized conversation starters based on client context
4. Maintain professional, clear, concise language
5. Include confidence scores for synthesized insights
6. Cite data sources for all statements
7. Flag missing or stale data explicitly

OUTPUT FORMAT: Return valid JSON matching the schema provided.

CONSTRAINTS:

- Never hallucinate data not present in inputs
- If data is missing, mark section as "Data unavailable"
- Keep each section concise (max 150 words)
- Prioritize recent information over historical
- Use bullet points for readability

User Prompt Template:

Generate a client briefing for the following meeting:

MEETING DETAILS:

- Client: {{clientName}}
- Meeting Time: {{meetingTime}}
- Meeting Type: {{meetingType}}
- Agenda: {{meetingAgenda}}

CLIENT DATA:

{{clientDataJSON}}

PORTFOLIO DATA:

{{portfolioDataJSON}}

RECENT COMMUNICATIONS:

{{emailSummaryJSON}}

MARKET DATA:

{{newsDataJSON}}

Generate briefing sections:

1. Client Overview
2. Recent Activity
3. Portfolio Snapshot
4. Alerts & Opportunities
5. Conversation Starters (2-3 items)
6. Meeting Context

Return JSON response matching this schema:

{{responseSchemaJSON}}

6.3 Response Schema

```
{
  "clientOverview": {
    "summary": "string",
    "keyFacts": ["string"],
    "dataSources": ["string"]
  },
  "recentActivity": {
    "lastMeetingDate": "ISO8601",
    "lastMeetingSummary": "string",
```



```

    "recentTransactions": [
      {
        "date": "ISO8601",
        "description": "string",
        "amount": "number"
      }
    ],
    "recentCommunications": "string"
  },
  "portfolioSnapshot": {
    "totalAUM": "number",
    "ytdPerformance": "number",
    "ytdBenchmark": "number",
    "assetAllocation": {
      "equities": "number",
      "fixedIncome": "number",
      "alternatives": "number",
      "cash": "number"
    },
    "topHoldings": [
      {
        "symbol": "string",
        "name": "string",
        "percentOfPortfolio": "number"
      }
    ]
  },
  "alertsAndOpportunities": [
    {
      "type": "warning|opportunity|info",
      "title": "string",
      "description": "string",
      "actionable": "boolean",
      "priority": "high|medium|low",
      "confidence": "high|medium|low"
    }
  ],
  "conversationStarters": [
    {
      "topic": "string",
      "rationale": "string"
    }
  ],
  "meetingContext": {
    "agendaSummary": "string",
    "outstandingActionItems": ["string"],
    "relevantNotes": ["string"]
  },
  "metadata": {
    "generatedAt": "ISO8601",
    "dataFreshness": {
      "crmData": "ISO8601",
      "portfolioData": "ISO8601",
      "emailData": "ISO8601",
      "newsData": "ISO8601"
    },
    "model": "string",
    "processingTimeMs": "number"
  }
}

```

6.4 Quality Assurance & Validation

Pre-Production Testing:

- Human evaluation: 100 sample briefings reviewed by advisors (target: 90% quality rating)
- A/B testing: Compare AI-generated vs. human-curated briefings (target: no significant quality difference)
- Edge case testing: Missing data, conflicting data, unusual client situations

Production Monitoring:

- Advisor feedback scores (track trend, investigate drops)
- Accuracy spot-checks: Random sample review (10 briefings/week)
- Hallucination detection: Flag unattributed statements, verify against source data
- Latency tracking: Alert if generation time > 10 seconds

Continuous Improvement:

- Prompt iteration based on feedback
- Fine-tuning potential: Collect high-quality examples for future model customization
- A/B testing of prompt variations

6.5 Guardrails & Safety

Content Filtering:

- Input sanitization: Remove PII before sending to external LLM APIs (if required by policy)
- Output validation: Check for inappropriate content, bias, or errors
- Fallback logic: If LLM returns invalid format, regenerate or show error

Bias Mitigation:

- Prompt explicitly instructs objective, unbiased language
- Monitor for demographic bias in conversation starters
- Human review process for bias detection

Hallucination Prevention:

- Prompt emphasizes "never fabricate data"
- Post-processing validates all statements against source data
- Confidence scoring surfaces uncertain synthesized information

7. Testing Strategy

7.1 Test Levels

Unit Testing:

- Target: 80%+ code coverage
- Frameworks: Jest (JavaScript), Pytest (Python)
- Focus: Individual functions, data transformations, utility methods

Integration Testing:

- Test API integrations with mocked external services
- Validate data flow between microservices
- Test message queue pub/sub patterns

End-to-End Testing:

- Simulate full user journeys (installation → briefing generation → feedback)
- Test across Windows and macOS environments
- Frameworks: Playwright (UI automation), Postman (API workflows)

Load Testing:

- Simulate 500 concurrent users generating briefings
- Test peak load scenarios (Monday morning, market open)
- Tools: Apache JMeter, k6

Security Testing:

- Automated vulnerability scanning (Snyk, OWASP ZAP)
- Penetration testing by external security firm (pre-launch)
- Regular security audits (quarterly)

7.2 Test Scenarios (Key User Flows)

Scenario 1: Happy Path Briefing Generation

- Given: Advisor has meeting in 30 minutes, all systems operational
- When: Calendar Service detects meeting
- Then: Briefing generates within 5 seconds, displays in widget, all sections populated

Scenario 2: Partial Data Availability

- Given: CRM data available, Portfolio API down
- When: Briefing generation triggered
- Then: Briefing shows CRM data, flags portfolio section as unavailable, completes within 5 seconds

Scenario 3: Network Failure During Meeting

- Given: Advisor viewing briefing, network connection lost
- When: Advisor interacts with widget
- Then: Cached briefing remains accessible, network-dependent actions disabled gracefully

Scenario 4: Concurrent Briefing Requests

- Given: Advisor has back-to-back meetings (3 in 1 hour)
- When: Multiple briefings triggered simultaneously
- Then: All briefings generate successfully, no data conflicts, proper queueing

Scenario 5: First-Time User Onboarding

- Given: New advisor with widget installed
- When: Advisor logs in for first time
- Then: Brief tutorial tooltip appears, first briefing auto-generated for next meeting

7.3 Acceptance Criteria for MVP Launch

Functional Acceptance:

- ✓ Calendar detection: 95%+ meetings detected correctly
- ✓ Briefing generation success rate: 99%+
- ✓ Data accuracy: 95%+ validated by advisor feedback
- ✓ Response time: 95th percentile $\leq$ 5 seconds

User Acceptance:

- ✓ User satisfaction: CSAT $\geq$ 4.2/5.0 (from 50+ survey responses)
- ✓ Adoption: 80%+ of target users actively using widget within 30 days
- ✓ Usability: Users can generate and view briefing with $\leq$ 3 clicks

Operational Acceptance:

- ✓ System uptime: 99%+ during pilot (30 days)
- ✓ Zero critical security incidents
- ✓ Support ticket volume: $\leq$ 0.1 tickets per user per month

Business Acceptance:

- ✓ Time savings: Validated 10+ minute reduction in prep time (from advisor surveys)
- ✓ Positive feedback: 80%+ of advisors rate briefings as "helpful" or "very helpful"
- ✓ Executive sponsor approval to proceed to full rollout

8. Deployment & Rollout Plan

8.1 Phased Rollout Strategy

Phase 0: Internal Alpha (Weeks 1-2)

- Audience: Product team + 5 friendly advisors
- Goal: Validate core functionality, identify critical bugs
- Success Criteria: Zero critical bugs, basic functionality working

Phase 1: Limited Pilot (Weeks 3-6)

- Audience: 25 advisors (mix of tech-savvy and typical users)
- Goal: User acceptance testing, gather feedback, refine UX
- Success Criteria: 80% CSAT, identify top 10 improvement areas

Phase 2: Expanded Pilot (Weeks 7-12)

- Audience: 100 advisors (representative sample across regions/teams)
- Goal: Validate scalability, performance under real load, training materials
- Success Criteria: Meet all launch acceptance criteria (Section 7.3)

Phase 3: General Availability (Week 13+)

- Audience: All eligible advisors (500+ users)
- Rollout: 100 users/week to monitor for issues
- Goal: Full adoption, continuous improvement based on feedback

8.2 Deployment Architecture

Widget Distribution:

- Enterprise software deployment tool (e.g., SCCM, Jamf for macOS)
- Auto-update mechanism: Check for updates daily, silent install minor versions
- Manual approval required for major version updates

Backend Deployment:

- Blue-green deployment strategy for zero-downtime updates
- Canary releases: 10% traffic → 50% → 100% over 24 hours
- Automated rollback on error rate spike or performance degradation

Configuration Management:

- Feature flags for gradual rollout of new capabilities
- Remote configuration updates without widget reinstall
- A/B testing framework for UX experiments

8.3 Training & Change Management

Training Materials:

- 5-minute product overview video
- Interactive tutorial (first-run experience)
- Knowledge base articles (searchable)
- Office hours: Weekly Q&A sessions (first 8 weeks)

Communication Plan:

- T-30 days: Executive announcement, value proposition
- T-14 days: Invitation to training webinar
- T-7 days: Installation instructions, support resources
- T-0 (Launch): Launch email with quick start guide
- T+7 days: Check-in survey, gather early feedback
- T+30 days: Success stories, adoption metrics shared

Change Champions:

- Identify 10-15 influential advisors as early adopters
- Provide exclusive early access and direct support
- Empower champions to advocate and support peers

8.4 Support Model

Tier 1 Support:

- Internal help desk (existing IT support team)
- Handle: Installation issues, login problems, basic troubleshooting
- Response SLA: 4 hours during business hours

Tier 2 Support:

- Product team (dedicated support engineer)
- Handle: Functional issues, data problems, complex troubleshooting
- Response SLA: Same business day
- Escalation path to engineering team

Tier 3 Support:

- Engineering team
- Handle: Bugs, system outages, architecture issues
- Response SLA: Immediate for critical issues (P0/P1)

Support Channels:

- Email: support@[company].com
- Slack channel: #wealth-companion-support (internal)
- Knowledge base: Searchable help articles
- In-app feedback: Submit feedback/bug reports directly from widget

9. Metrics & Analytics

9.1 Product Metrics Dashboard

Usage Metrics:

- Daily Active Users (DAU), Weekly Active Users (WAU)
- Briefings generated per user per day (average, median, 90th percentile)
- Session duration (time spent viewing briefing)
- Section engagement (% of users expanding each section)
- Feedback submission rate and sentiment

Performance Metrics:

- Briefing generation time (P50, P95, P99)
- API response times for each integration
- Error rate by component
- System uptime (target: 99.5%)

Business Impact Metrics:

- Time saved per advisor (surveyed monthly)
- Advisor satisfaction (CSAT, NPS)
- Meeting quality improvement (client feedback)
- Opportunity conversion rate (opportunities identified → actions taken)

9.2 Data Collection & Privacy

Telemetry:

- Anonymous usage data collected (no PII)
- Opt-out available for advisors (with notification to manager)
- Data aggregated for reporting (no individual tracking)

Event Tracking:

- Widget opened/closed
- Briefing viewed
- Section expanded/collapsed
- Feedback submitted (rating + optional comment)
- Follow-up task created
- Error encountered

Data Retention:

- Usage metrics: 2 years
- Error logs: 90 days (longer if under investigation)
- Audit logs: 7 years (compliance requirement)

9.3 Success Criteria (6-Month Post-Launch)

Metric	Target	Measurement Method
Adoption Rate	90%+ of eligible advisors	Active users (≥ 3 briefings/week)
Time Savings	50+ min/advisor/week	Monthly survey (validated sample)

Metric	Target	Measurement Method
User Satisfaction	CSAT ≥ 4.3/5.0	In-app feedback + quarterly survey
Briefing Accuracy	95%+ advisor validation	Random spot-checks + feedback
System Reliability	99.5%+ uptime	Automated monitoring
Business ROI	5x cost of implementation	Time savings × avg advisor comp

10. Risks & Mitigation

10.1 Technical Risks

Risk 1: LLM API Instability

- **Likelihood:** Medium
- **Impact:** High (briefings fail to generate)
- **Mitigation:**
 - Multi-provider strategy (GPT-4o + Claude fallback)
 - Circuit breaker to prevent cascading failures
 - Cached fallback: Show previous briefing with staleness warning
 - SLA monitoring and automatic escalation

Risk 2: Integration Failures (CRM, Portfolio systems)

- **Likelihood:** Medium
- **Impact:** High (incomplete or missing data)
- **Mitigation:**
 - Graceful degradation: Show available data, flag missing sections
 - Retry logic with exponential backoff
 - Dedicated monitoring and alerts for integration health
 - Fallback to direct database queries (read-only replicas)

Risk 3: Performance Degradation Under Load

- **Likelihood:** Low-Medium
- **Impact:** Medium (slow briefing generation, poor UX)
- **Mitigation:**
 - Load testing before launch (simulate 2x expected peak load)
 - Auto-scaling infrastructure (horizontal scaling)
 - Rate limiting and request queuing
 - Performance monitoring and alerting (P95 latency > 8 sec)

Risk 4: Data Quality Issues

- **Likelihood:** Medium
- **Impact:** Medium (inaccurate briefings, loss of trust)
- **Mitigation:**
 - Data validation layer before AI processing
 - Source attribution for all data points (advisor can verify)
 - Confidence scoring for synthesized insights
 - Feedback loop: Advisors flag inaccuracies, engineering investigates

10.2 Business & Adoption Risks

Risk 5: Low User Adoption

- **Likelihood:** Low-Medium
- **Impact:** High (product failure)
- **Mitigation:**
 - Executive sponsorship and top-down communication
 - Change champions program (peer advocates)
 - Demonstrate clear value in training (time savings, better meetings)
 - Frictionless onboarding (auto-install, zero configuration)
 - Gamification: Leaderboard of widget users (opt-in)

Risk 6: Resistance from Advisors ("AI will replace me")

- **Likelihood:** Low
- **Impact:** Medium (adoption resistance, morale)
- **Mitigation:**
 - Clear messaging: Widget augments advisors, doesn't replace
 - Emphasize time savings for high-value activities (relationship building)
 - Advisory board: Include advisors in product development
 - Highlight success stories early and often

Risk 7: Insufficient ROI / Value Perception

- **Likelihood:** Low
- **Impact:** High (product cancellation)
- **Mitigation:**
 - Focus MVP on high-value, high-frequency use case (pre-meeting prep)
 - Measure and communicate time savings explicitly
 - Gather testimonials from pilot users
 - Track leading indicators: usage frequency, feedback sentiment

10.3 Security & Compliance Risks

Risk 8: Data Breach or Unauthorized Access

- **Likelihood:** Low
- **Impact:** Critical (regulatory fines, reputational damage)
- **Mitigation:**
 - Zero persistent client data storage in widget
 - End-to-end encryption (TLS 1.3, AES-256)
 - RBAC with least-privilege principle
 - Regular security audits and penetration testing
 - Incident response plan (pre-defined, tested)

Risk 9: Regulatory Non-Compliance

- **Likelihood:** Low
- **Impact:** High (fines, operational restrictions)
- **Mitigation:**
 - Compliance review before launch (legal, risk, compliance teams)
 - Audit logging (tamper-proof, 7-year retention)

- Data residency compliance (process data in approved regions)
- Regular compliance training for product team
- Third-party compliance audit (annual)

Risk 10: AI Hallucinations or Errors

- **Likelihood:** Medium
- **Impact:** Medium-High (incorrect advice, client harm)
- **Mitigation:**
 - Strict prompt engineering ("never fabricate data")
 - Post-processing validation (check against source data)
 - Confidence scoring surfaced to advisors
 - Human-in-the-loop: Advisors review before acting on insights
 - Feedback mechanism to flag errors
 - Monthly quality review: Random sample of briefings audited

10.4 Operational Risks

Risk 11: Dependency on Third-Party APIs (LLM providers)

- **Likelihood:** Medium
- **Impact:** Medium (service disruption)
- **Mitigation:**
 - Multi-provider contracts (OpenAI + Anthropic)
 - Monitor provider status pages, auto-failover
 - Explore local model hosting for mission-critical use cases (future)
 - Communicate transparently with users during outages

Risk 12: Insufficient Engineering Resources

- **Likelihood:** Low
- **Impact:** High (delayed launch, poor quality)
- **Mitigation:**
 - Dedicated cross-functional team (5-8 engineers, 1 PM, 1 designer)
 - Clear prioritization and scope management
 - Agile methodology with 2-week sprints
 - Buffer in timeline for unexpected issues (20%)

11. Open Questions & Decisions Needed

11.1 Technical Decisions

Q1: Hosting Strategy

- **Question:** Should backend services run in cloud (AWS/Azure) or on-premises?
- **Options:**
 - A) Cloud-native (AWS): Faster deployment, managed services, easier scaling
 - B) On-premises: Greater control, data residency, compliance comfort
 - C) Hybrid: Widget + API gateway on-prem, AI services in cloud
- **Decision Owner:** CTO, Head of Infrastructure
- **Timeline:** Decide by Week 2 of project kickoff

Q2: LLM Provider Strategy

- **Question:** Should we build with ability to fine-tune models or rely on off-the-shelf APIs?
- **Options:**
 - A) API-only (faster MVP, lower initial cost)
 - B) Plan for fine-tuning (higher quality long-term, data collection)
- **Decision Owner:** Head of AI/ML, Product Management
- **Timeline:** Decide by Week 4

Q3: Real-Time Updates

- **Question:** Should briefings auto-update if data changes before meeting?
- **Options:**
 - A) Static: Generate once, no updates (simpler)
 - B) Dynamic: Re-generate if significant data change (better accuracy)
- **Decision Owner:** Product Management, Engineering Lead
- **Timeline:** Decide during design phase (Week 6)

11.2 Business Decisions

Q4: Pricing/Funding Model

- **Question:** How should this product be funded and potentially charged?
- **Options:**
 - A) Internal project: No chargeback, full budget from IT/Innovation
 - B) Cost allocation: Charge business units per user
 - C) Freemium: Basic free, charge for advanced features
- **Decision Owner:** CFO, Business Unit Heads
- **Timeline:** Decide before pilot begins (Week 8)

Q5: Scope Expansion Priority

- **Question:** After MVP, which use case should be prioritized next?
- **Options:**
 - A) Prospect research automation
 - B) Portfolio health monitoring with alerts
 - C) Regulatory change tracking
- **Decision Owner:** Product Management, Advisory Board
- **Timeline:** Decide during pilot (Week 10), inform roadmap

11.3 Compliance & Governance

Q6: Client Disclosure Requirements

- **Question:** Do clients need to be notified that AI is used in advisor prep?
- **Options:**
 - A) Yes, explicit disclosure required (conservative)
 - B) No, internal tool, no client-facing output (pragmatic)
- **Decision Owner:** Chief Compliance Officer, Legal
- **Timeline:** Decide before pilot (Week 6)

Q7: Data Retention Policy

- **Question:** How long should briefings be retained (if at all)?

- **Options:**
 - A) Ephemeral: Delete after meeting (privacy-first)
 - B) 90 days: Allow historical access (advisor convenience)
 - C) 7 years: Match document retention policy (compliance)
- **Decision Owner:** Compliance, Legal, Privacy Officer
- **Timeline:** Decide before launch (Week 10)

12. Assumptions & Dependencies

12.1 Assumptions

Technical Assumptions:

- OpenAI and Anthropic APIs will remain available and stable
- Internal CRM and Portfolio systems have documented APIs with sufficient data access
- Advisors have modern laptops (Windows 10+, macOS 12+) with admin rights for software installation
- Network bandwidth is sufficient (100 Mbps recommended)
- Calendar integration can be achieved via Microsoft Graph or Google Calendar APIs

Business Assumptions:

- Advisors spend 10-15 min on manual pre-meeting prep currently (to be validated)
- Time savings will be valued by advisors and translate to better client service
- Executive sponsorship will remain strong throughout project
- Pilot users will be representative of broader population
- Budget allocated (\$500K-\$1M for MVP) will be sufficient

User Assumptions:

- Advisors have 3-5 client meetings per day on average
- Advisors are comfortable with desktop software (not just web apps)
- Advisors trust AI-generated insights if source-attributed and accurate
- Feedback will be provided honestly to improve product

12.2 Dependencies

Internal Dependencies:

- **CRM Team:** API access granted, documentation provided, support for integration troubleshooting
- **Portfolio Management Team:** API/database access, data dictionary, performance SLA commitment
- **IT Security Team:** Security review completed, approval for deployment, SSO integration
- **Compliance Team:** Product approval, ongoing audit support, regulatory guidance
- **Change Management Team:** Communication templates, training resources, rollout coordination
- **Executive Sponsor:** Budget approval, organizational advocacy, executive steering committee participation

External Dependencies:

- **OpenAI / Anthropic:** API access, rate limit increases for production, support for critical issues
- **Microsoft / Google:** Calendar API access, webhook reliability, support for OAuth issues
- **Cloud Provider (AWS/Azure):** Infrastructure availability, service uptime SLAs, cost predictability
- **Desktop Deployment Tool Vendor (SCCM/Jamf):** Support for Electron app deployment, troubleshooting

Timeline Dependencies:

- CRM API access must be available by Week 3 (critical path)

- Security review must complete by Week 8 to allow pilot start on Week 10
- Executive sponsor approval required by Week 6 for pilot resourcing
- Pilot advisor recruitment must begin Week 6 for Week 10 start

13. Appendices

Appendix A: Glossary

Term	Definition
AUM	Assets Under Management - total market value of client assets managed by advisor
CSAT	Customer Satisfaction Score - rating scale (typically 1-5) measuring user satisfaction
CRM	Customer Relationship Management - system for managing client data and interactions
HNW	High Net Worth - individuals with investable assets \$1M-\$5M
LLM	Large Language Model - AI model trained on text data (e.g., GPT-4, Claude)
MVP	Minimum Viable Product - first version with core features to validate concept
NPS	Net Promoter Score - metric measuring user loyalty (-100 to +100 scale)
RBAC	Role-Based Access Control - security model restricting access based on user roles
SSO	Single Sign-On - authentication allowing one login for multiple systems
UHNW	Ultra High Net Worth - individuals with investable assets > \$30M

Appendix B: References

[^74] Atlassian - How to create a product requirements document (PRD)

[^75] Product School - The Only Product Requirements Document (PRD) Template

[^79] Product Compass - A Proven AI PRD Template by Miqdad Jaffer

[^80] CMARIX - Fintech in Wealth Management: Shaping Finance Industry

[^83] Biz4Group - Fintech in Wealth Management: The Future Is Now

[^86] KFinTech - Benefits of FinTech in Wealth Management

[^94] Salesforce - Wealth Management Software for Financial Advisors

[^96] SS&C Black Diamond - Top 10 Automated Workflows for Financial Advisors

[^98] ForgeXRM - 5 Key Challenges Solved by CRM for Wealth Management Firms

[^99] Alation - AI in Financial Services: Benefits, Applications & Implementation

[^102] AWS Blog - Agentic AI in Financial Services

[^111] Blueflame AI - AI Roadmap Guide for Mid-Size Financial Services Firms

Appendix C: Revision History

Version	Date	Author	Changes
0.1	Oct 15, 2025	Product Team	Initial draft outline
0.5	Oct 17, 2025	Product Team	Added technical architecture, AI specs
1.0	Oct 18, 2025	Product Team	Complete first draft for stakeholder review

END OF DOCUMENT

This PRD is a living document and will be updated iteratively based on stakeholder feedback, technical discoveries, and evolving requirements. Please direct questions or suggestions to the Product Management team.



1. <https://www.cmarix.com/blog/fintech-in-wealth-management/>
2. <https://radixweb.com/blog/desktop-application-development-guide>
3. <https://learn.microsoft.com/en-us/dotnet/architecture/modern-web-apps-azure/common-web-application-architectures>
4. <https://www.atlassian.com/agile/product-management/requirements>
5. <https://www.biz4group.com/blog/fintech-in-wealth-management>
6. <https://www.aha.io/roadmapping/guide/requirements-management/what-is-a-good-product-requirements-document-template>
7. <https://www.productleadership.com/blog/product-requirements-document-using-genai/>
8. <https://www.kfintech.com/benefits-of-fintech-in-wealth-management-explained/>
9. <https://www.atlassian.com/software/confluence/templates/product-requirements>
10. <https://chatprd.ai/resources/using-ai-to-write-prd>
11. <https://www.synechron.com/en-in/insight/what-fintech-product-managers-need-know-2024>
12. <https://www.news.aakashg.com/p/product-requirements-documents-prds>
13. <https://www.news.aakashg.com/p/ai-prd>
14. <https://hellopmp.co/fintech-product-management-in-india/>
15. <https://productschool.com/blog/product-strategy/product-template-requirements-document-prd>
16. <https://www.chatprd.ai/templates>
17. <https://www.salesforce.com/in/financial-services/wealth-management-software/>
18. <https://www.wealthbox.com>
19. <https://www.sscblackdiamond.com/discover/blog/top-10-automated-workflows-every-financial-advisor-should-build-in-their-crm/>
20. <https://dynatechconsultancy.com/blog/best-crm-software-for-financial-advisors-wealth-management>
21. <https://www.crmsoftwareblog.com/2025/02/5-key-challenges-solved-by-crm-for-wealth-management-firms/>
22. <https://www.alation.com/blog/ai-financial-services-benefits-implementation/>
23. <https://sam-solutions.com/blog/how-to-make-a-desktop-application/>
24. <https://formslogic.com/the-best-crm-for-financial-advisors-choosing-the-right-fit-for-your-workflow/>
25. <https://aws.amazon.com/blogs/awsmarketplace/agent-ai-solutions-in-financial-services/>
26. <https://www.notion.com/templates/category/product-requirements-doc>
27. <https://gototags.com/desktop-app/system-requirements>
28. <https://www.engagebay.com/blog/crm-for-financial-advisors/>
29. <https://www.360factors.com/blog/six-steps-implement-artificial-intelligence-in-banking/>
30. <https://www.equisoft.com/products/advisor-crm-connect>
31. <https://www.bis.org/fsi/publ/insights63.pdf>
32. <https://www.advisorengine.com/action-magazine/articles/how-advisors-use-crm-to-unclutter-their-client-management>
33. <https://www.blueflame.ai/blog/ai-roadmap-guide-for-financial-services>
34. <https://community.devexpress.com/blogs/news/archive/2024/02/09/modern-desktop-apps-and-their-complex-architectures.aspx>
35. <https://www.hustlebadger.com/what-do-product-teams-do/prd-template-examples/>
36. <https://www.assetmark.com/blog/financial-advisor-workflows-best-practices-for-optimal-efficiency>
37. <https://www.figma.com/templates/prd-template/>
38. <https://www.productcompass.pm/p/ai-prd-template>
39. <https://www.productplan.com/glossary/product-requirements-document/>
40. <https://chisellabs.com/blog/how-to-write-prd-using-ai/>